



Hochschule Ravensburg-Weingarten

University of Applied Sciences

Lidar and Radar Systems

Task: Evaluation of an 3D Object Detector

Sr. No.	Student Name	Matriculation Numbers	Study Programm
1	Arjun Rameshbhai Beladiya	28744163	Mechatronics
2	Dhaval Jagdish Mistry	20644015	Mechatronics

Guided by:

Prof. Dr. rer. nat. Stefan Elser

Contents

1	Introduction	3
2	Methodology	3
2.1	Loading the KITTI-360 Dataset	3
2.2	YOLO Semantic Segmentation on Images	3
2.3	Projection of LiDAR Points onto the Image	3
2.3.1	Transformation from LiDAR to Camera Frame	4
2.3.2	Conversion Projection Matrix	4
2.3.3	Using P_rect_00 to Project 3D Points into 2D Image Plane	4
2.4	Matching Projected Points with YOLO Masks	5
2.5	Projecting Lidar points into 3D Space	5
2.6	Comparing LiDAR Points with Ground Truth Bounding Boxes	5
2.6.1	Precision and Recall	5
2.7	Flow chart	6
3	Results	7
4	Evaluation	17
4.1	Precision Threshold Justification	17
4.2	Image specific errors	17
4.2.1	Point Cloud Color Leakage into Multiple Boxes	17
4.2.2	Missing Ground Truth Bounding Boxes	17
4.2.3	Too many GT boxes on one car	17
5	Conclusion	18
6	References	18

1 Introduction

This task focuses on detecting cars in LiDAR points when they are projected on 2D image as well as mapped in 3D space. To perform this task, KITTI 360 sample (3) dataset was used; which is sub Dataset of KITTI 360(2). Further in this report, we will discuss Methodology, Results and their Evaluation.

2 Methodology

Before starting this project, we have to import the libraries open3d, cv2, numpy, os, and json to perform essential tasks such as 3D point cloud visualization, image processing, numerical computations, file management, and data visualization. The detailed approach to our solution is explained below.

2.1 Loading the KITTI-360 Dataset

In this task, The used KITTI-360 sample dataset (3) contains all necessary data such as camera images, LiDAR scans, calibration file, and additionally contains 3D ground truth bounding boxes. To utilize these data effectively, we first need to understand the content of these files.

- **data_2d_raw/**: Contains the raw camera data. Specifically, the `image_00/data_rect/` subfolder contains 20 images from the left camera.
- **data_3d_raw/**: Contains raw 3D LiDAR point cloud data (`.bin` file) captured by veodyne sensor.
- **calibration/**: Includes all necessary calibration files such as `calib_cam_to_velo`, `perspective`.
- **bboxes_3D_cam0/**: Contains the annotated ground truth 3D bounding boxes in JSON format.

2.2 YOLO Semantic Segmentation on Images

First we need to mask the cars in 2D image which is later used for finding LiDAR points corresponding to cars. For this semantic segmentation, we have chose the YOLOv8x-seg (1) model because it provides a good balance between performance and accuracy, producing precise pixel-wise masks. For this task, we will only focus on cars. Therefore, all other objects (Pedestrians, cyclists, Truck, Traffic Signs, etc) are filtered out. we have used different colors for different segmentated cars.

2.3 Projection of LiDAR Points onto the Image

Now to accurately project 3D LiDAR points onto 2D camera images, first we need to align the coordinate systems of the LiDAR sensor and the camera. which we have done with the help of `calib_cam_to_velo` and `perspective` file given KITTI 360 subdataset (3).

2.3.1 Transformation from LiDAR to Camera Frame

`calib_cam_to_velo` file contains 12 numerical values of 3×4 transformation matrix which maps coordinates from the camera frame to the LiDAR frame. but we want to project LiDAR points into the camera frame, so we took inverse of the transformation matrix to convert LiDAR coordinates into the camera coordinate system:

$$T_{\text{velo} \rightarrow \text{cam}} = (T_{\text{cam} \rightarrow \text{velo}})^{-1}$$

Now we have all 3D points from the LiDAR's world into the same spatial frame as camera.

2.3.2 Conversion Projection Matrix

The LiDAR sensor captures points in its local 3D coordinate frame, which is represented as:

$$\mathbf{P}_{\text{lidar}} = [x \ y \ z]^T$$

To apply transformation, we have to convert 3D points to homogeneous coordinates by adding an extra 1 to make it a 4×1 vector.

$$\mathbf{P}_{\text{hom}} = [x \ y \ z \ 1]^T$$

Now, we can do both rotation and translation and by multiplying transformation matrix to \mathbf{P}_{hom} we can convert points in the camera frame:

$$\mathbf{P}_{\text{cam}} = T_{\text{velo} \rightarrow \text{cam}} \cdot \mathbf{P}_{\text{hom}}$$

2.3.3 Using `P_rect_00` to Project 3D Points into 2D Image Plane

Now, with the help of projection matrix `P_rect_00` we will project 3D points from the camera coordinate frame onto the 2D image plane, which is provided in the `perspective.txt` calibration file. This matrix contains both the intrinsic camera parameters and rectification transformation to ensure the projected points align with the undistorted (rectified) image.

$$P_{\text{rect}} = \begin{bmatrix} f_x & 0 & c_x & t_x \\ 0 & f_y & c_y & t_y \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

Where:

- f_x, f_y are the focal lengths in pixels (scaling factors),
- c_x, c_y are the optical center coordinates (principal point),
- t_x, t_y are translation offsets (often 0 or adjusted for stereo calibration).

Now we have all required parameters to project each 3D points $\mathbf{P}_{\text{cam}} = [x \ y \ z \ 1]^T$ into image plane using this equation:

$$\begin{bmatrix} u \\ v \\ w \end{bmatrix} = P_{\text{rect}_00} \cdot \mathbf{P}_{\text{cam}}$$

In the next step we have applied perspective division to convert result into valid 2D pixel coordinates (u', v')

$$u' = \frac{u}{w}, \quad v' = \frac{v}{w}$$

These pixel coordinates represent where each 3D LiDAR point appears on the 2D image.

2.4 Matching Projected Points with YOLO Masks

Now, the LiDAR points are being projected onto the image plane. But which point is corresponds to which car is still remains a problem. To solve this pixel-wise segmentation masks generated by the YOLOv8x-seg model (1) is used. Every detected car has a binary mask, which tells us which specific pixel belongs to that specific car. These masks are stored as 2D arrays with values of 1 indicating the presence of the object and 0 elsewhere. Now, if the pixel value at (u', v') is 1, then the LiDAR point is inside the corresponding car region and to differentiate every car points. A unique RGB have been assigned color to every LiDAR point laying inside every car. If the point is outside the car, then it considered background and is colored black.

2.5 Projecting Lidar points into 3D Space

To visualize the LiDAR points in 3D space, we have used the Open3D library. The projected points are represented as a point cloud, which can be visualized in a 3D coordinate system. This point cloud was created by taking the projected 2D pixel coordinates and mapping them back to their corresponding 3D coordinates using the inverse of the projection matrix. After that we have reapplied colors(which was saved during Yolo segmentation) to the points corresponding to cars.

2.6 Comparing LiDAR Points with Ground Truth Bounding Boxes

To evaluate the accuracy of the LiDAR points against the ground truth bounding boxes, we have used the 3D bounding box annotations provided in the KITTI 360 Sample Dataset (3). The ground truth bounding boxes are defined in the 3D space of LiDAR and contain the 3D coordinates of the corners of each box. To evaluate the projected LiDAR points with the ground truth bounding boxes, first we have checked if the projected points fall within the boundaries of these boxes.

2.6.1 Precision and Recall

To evaluate the performance of our LiDAR-image fusion and segmentation pipeline, we use two primary formulae: Precision and Recall

Precision: It measures the accuracy of the detections made by the model. It is the ratio of (TP) to (TP + FP). Formula:

$$\text{Precision} = \frac{\text{True Positives (TP)}}{\text{True Positives (TP)} + \text{False Positives (FP)}} \quad (1)$$

Recall: It quantifies the model's ability to detect all relevant objects in the scene. It is the ratio of (TP) to (TP + FN). Formula:

$$\text{Recall} = \frac{\text{True Positives (TP)}}{\text{True Positives (TP)} + \text{False Negatives (FN)}} \quad (2)$$

Where:

- **True Positive (TP):** A point correctly assigned to the car it belongs to, i.e., it lies within the predicted 3D GT bounding box and has the correct segmentation label.
- **False Positive (FP):** A point that is included in a predicted 3D bounding box but does not belong to that car.
- **False Negative (FN):** A point belonging to a car (based on segmentation) that was missed by the predicted 3D bounding box.

By the help of Recall and Precision values, we can evaluate the performance of the object detection model. The more we get the value close to 1, the better the performance. Below [1]i have attached the result of Precision and Recall of our images in the form of graph.

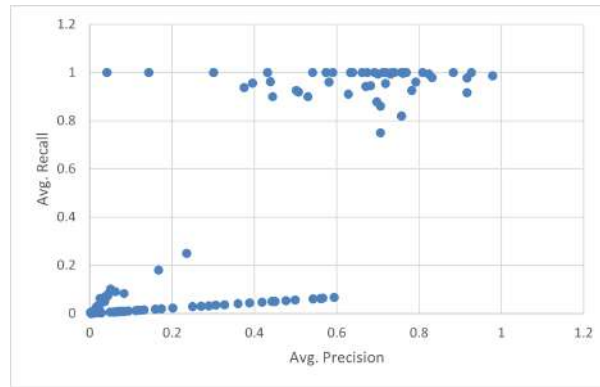


Figure 1: Avg. Precision vs Avg. Recall graph

2.7 Flow chart

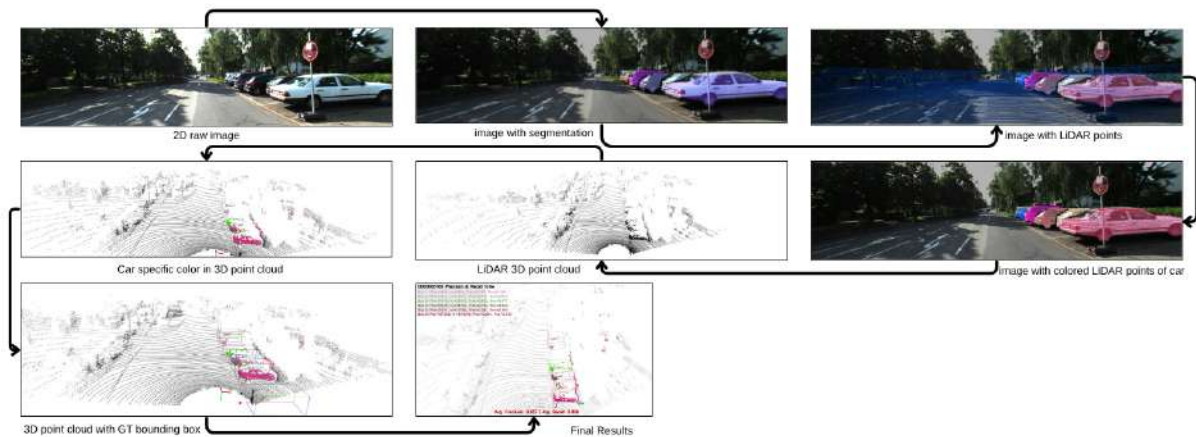


Figure 2: Flowchart

3 Results

Final result of images, which contains points belong to a car, points inside bounding box (TP), precision, recall, Avg. precision and Avg. recall values.



Figure 3: 2D View: 0000000100.png



Figure 4: 2D View: 0000000250.png

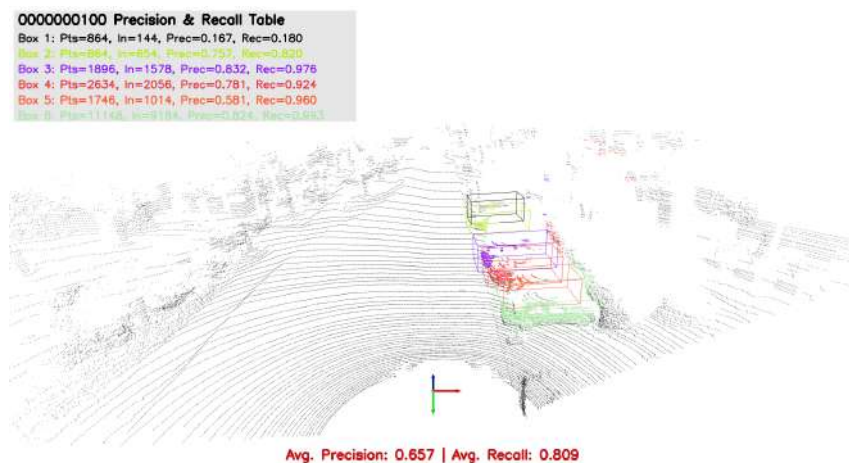


Figure 5: 3D Point cloud: 100.png

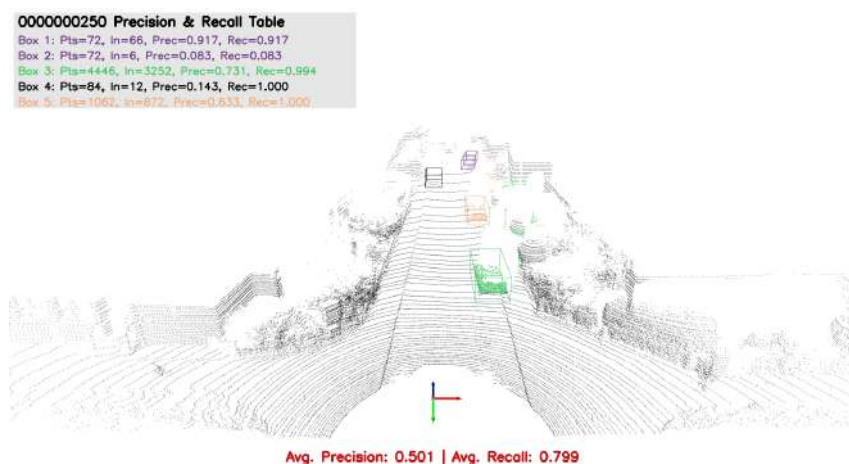


Figure 6: 3D Point cloud: 250.png



Figure 7: 2D View: 0000000360.png



Figure 8: 2D View: 0000000570.png

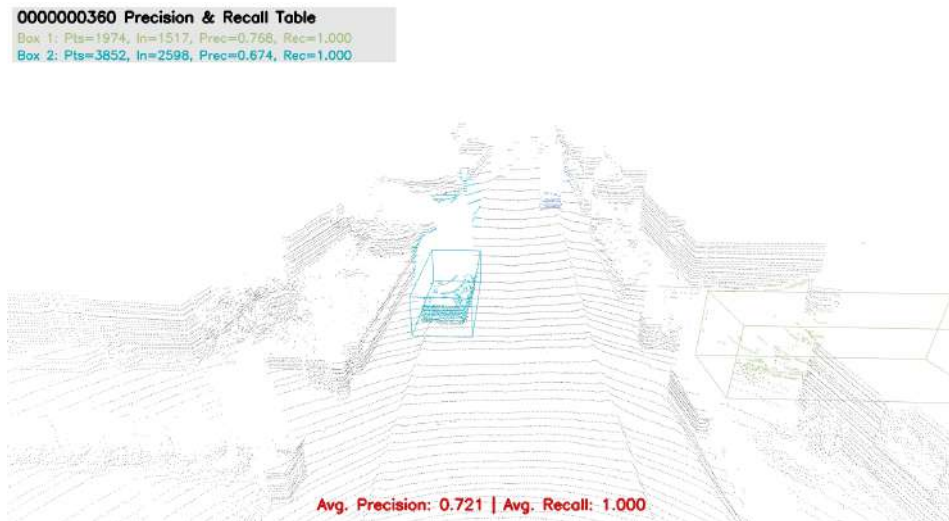


Figure 9: 3D Point cloud: 360.png




Figure 10: 3D Point cloud: 570.png



Figure 11: 2D View: 0000000644.png



Figure 12: 2D View: 0000000729.png



Figure 13: 3D Point cloud: 644.png

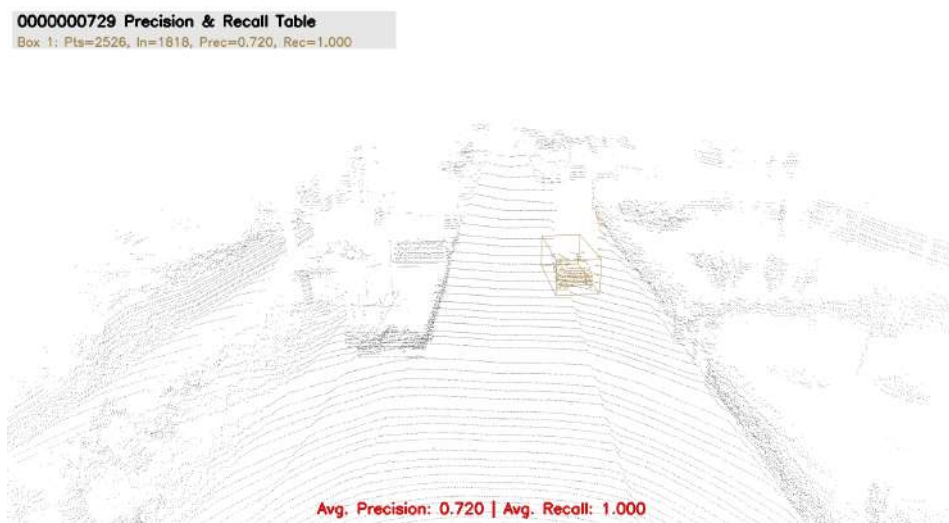


Figure 14: 3D Point cloud: 729.png




Figure 15: 2D View: 0000000947.png



Figure 16: 2D View: 0000001134.png

0000000947 Precision & Recall Table

Box 1:	Pts=222, In=168, Prec=0.757, Rec=1.000
Box 2:	Pts=13980, In=624, Prec=0.045, Rec=0.081
Box 3:	Pts=13980, In=7092, Prec=0.507, Rec=0.919
Box 4:	Pts=4626, In=78, Prec=0.017, Rec=0.029
Box 5:	Pts=4626, In=198, Prec=0.043, Rec=0.072




Avg. Precision: 0.317 | Avg. Recall: 0.500

Figure 17: 3D Point cloud: 947.png

0000001134 Precision & Recall Table

Box 1:	Pts=10734, In=7431, Prec=0.692, Rec=1.000
Box 2:	Pts=4050, In=2544, Prec=0.628, Rec=0.910
Box 3:	Pts=4050, In=252, Prec=0.062, Rec=0.090
Box 4:	Pts=9018, In=4884, Prec=0.542, Rec=1.000



Avg. Precision: 0.481 | Avg. Recall: 0.750

Figure 18: 3D Point cloud: 1134.png



Figure 19: 2D View: 0000001280.png



Figure 20: 2D View: 0000001338.png

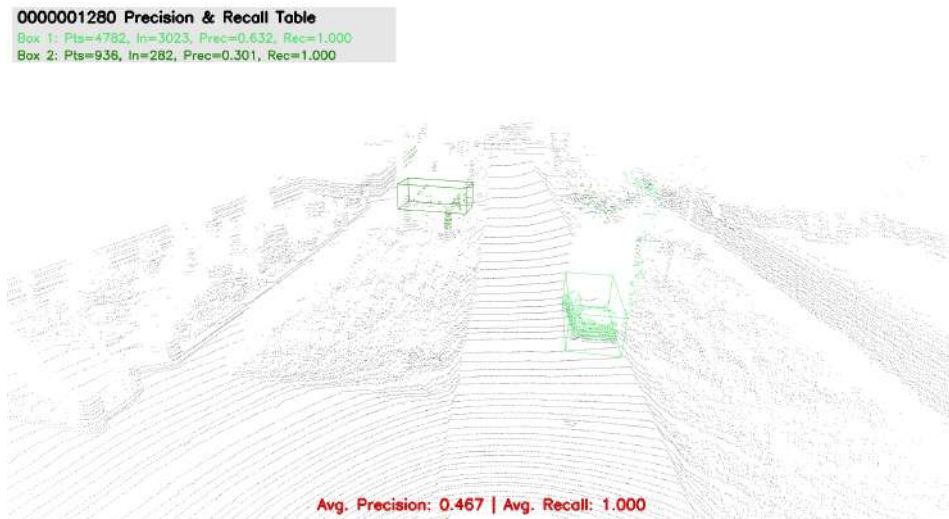


Figure 21: 3D Point cloud: 1280.png

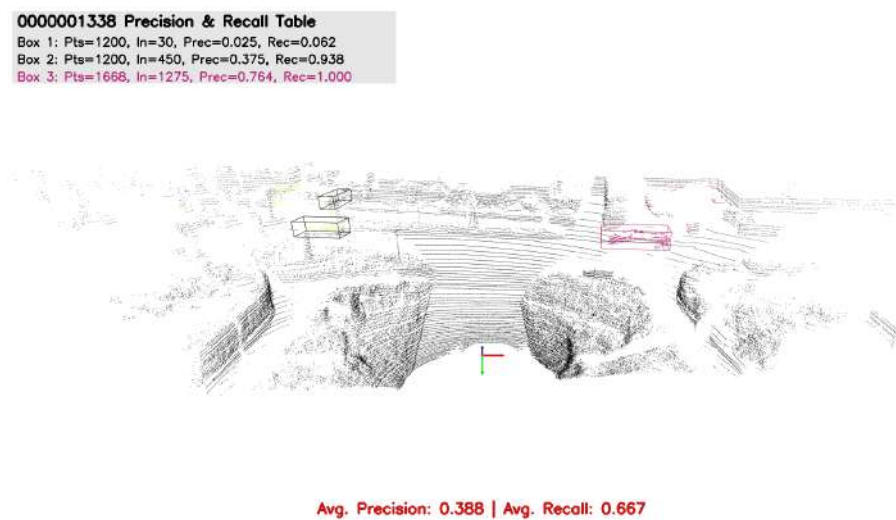


Figure 22: 3D Point cloud: 1338.png



Figure 23: 2D View: 0000001461.png



Figure 24: 2D View: 0000001597.png

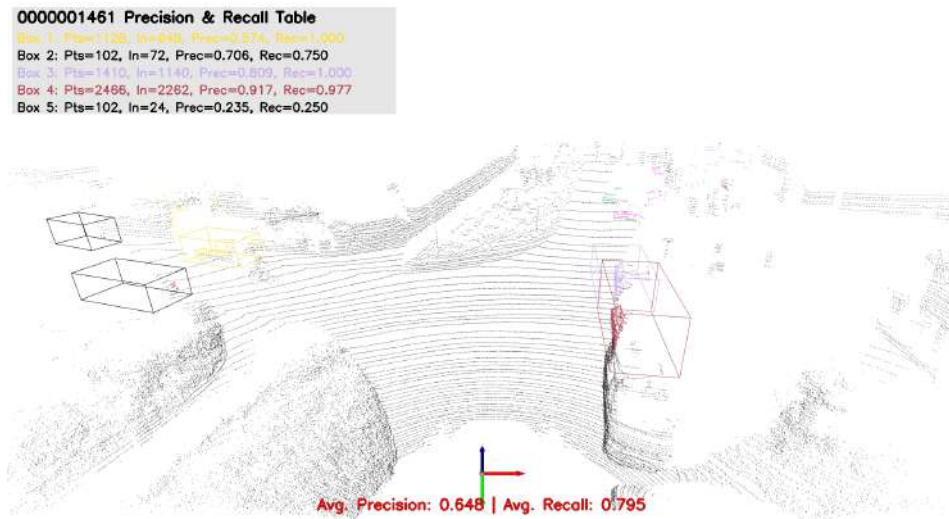


Figure 25: 3D Point cloud: 1461.png

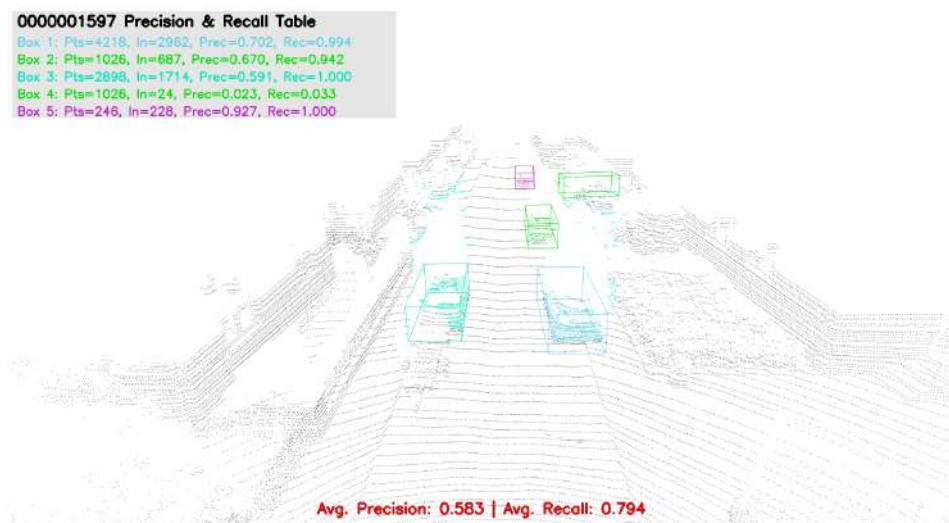


Figure 26: 3D Point cloud: 1597.png




Figure 27: 2D View: 0000001791.png



Figure 28: 2D View: 0000001866.png

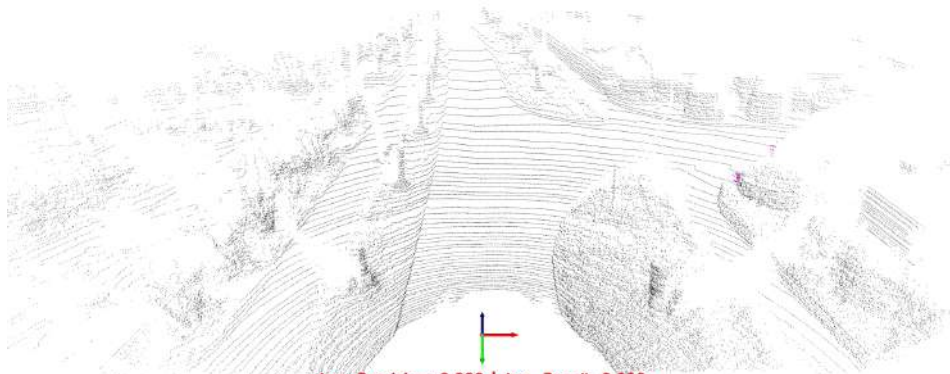
0000001791 Precision & Recall Table



Avg. Precision: 0.000 | Avg. Recall: 0.000

Figure 29: 3D Point cloud: 1791.png

0000001866 Precision & Recall Table



Avg. Precision: 0.000 | Avg. Recall: 0.000

Figure 30: 3D Point cloud: 1866.png



Figure 31: 2D View: 0000002033.png



Figure 32: 2D View: 0000002098.png

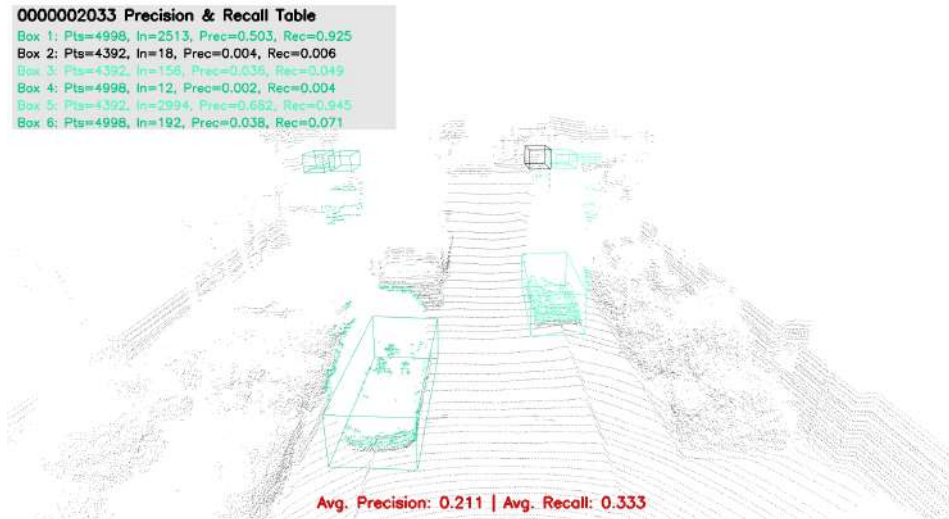


Figure 33: 3D Point cloud: 2033.png




Figure 34: 3D Point cloud: 2098.png



Figure 35: 2D View: 0000002449.png



Figure 36: 2D View: 0000002717.png

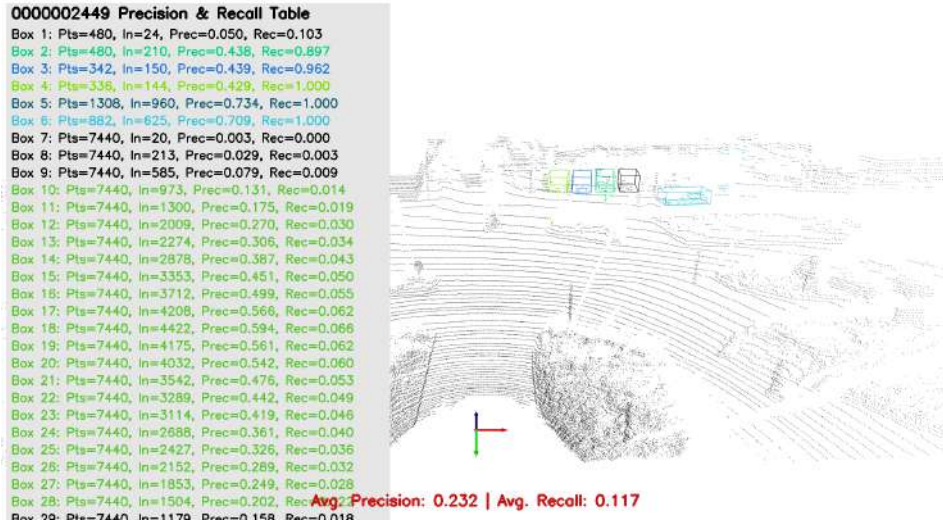


Figure 37: 3D Point cloud: 2449.png

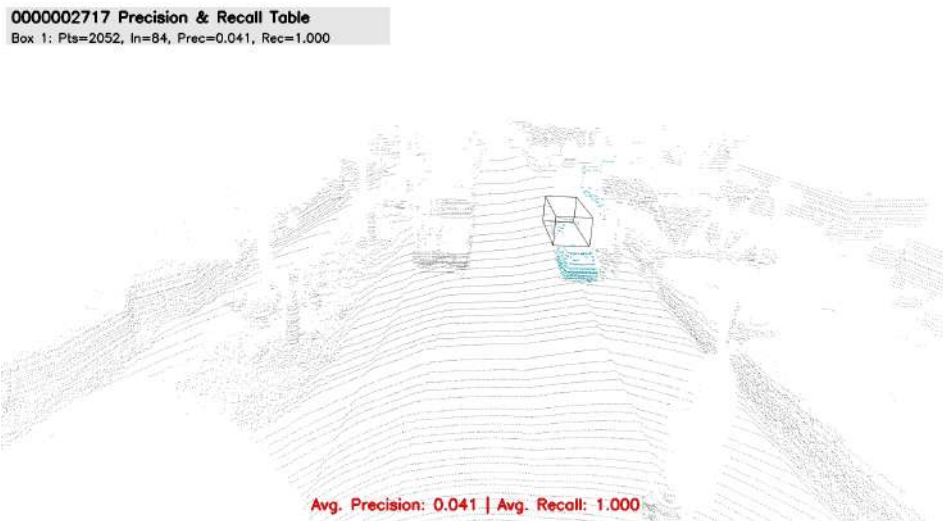


Figure 38: 3D Point cloud: 2717.png



Figure 39: 2D View: 0000002903.png



Figure 40: 2D View: 0000002939.png

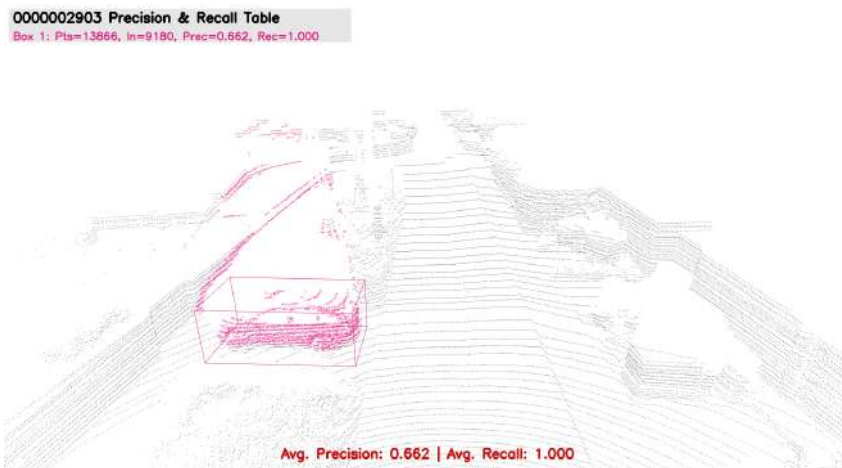


Figure 41: 3D Point cloud: 2903.png



Figure 42: 3D Point cloud: 2939.png

4 Evaluation

To evaluate the result, we will use the relative amount of points. Since point density depends on object distance and occlusion, absolute counts can be misleading. so precision and recall normalize these differences, providing a fairer quality measure. As a result we will evaluate this model based on the precision and recall values, obtained from the YOLOv8x-seg (1) model.

Detection Category	Precision	Count
Total number of detections		110
Total number of good detections	(Precision > 0.6)	34
Total number of poor detections	(0.6 > Precision > 0.35)	25
Total number of bad detections	(Precision < 0.35)	51

Table 1: Summary of detection evaluation results based on precision thresholds.

4.1 Precision Threshold Justification

If the value of precision is greater than 0.6, we considered it as a correct detection. This threshold ensures that the majority of the points inside the predicted 3D bounding box truly belong to the car, while allowing a small margin of error due to real-world factors such as noise, occlusions, or segmentation inaccuracies. A precision above 0.6 strikes a practical balance between strict accuracy and robustness, making the detection reliable enough for downstream tasks such as tracking or decision-making.

4.2 Image specific errors

4.2.1 Point Cloud Color Leakage into Multiple Boxes

In this scenes [Image Ids: 947(17), 1134(18), 1461(25), 1597(26), 2033(33), 2098(34), 2449(37)], the same color (i.e., segmented car) point cloud was partially spread across two or more bounding boxes. If the second bounding box contains only leaked points from the first car, it gets falsely counted as TP for both boxes, which leads to artificially inflated precision for incorrect boxes. This occurs because the segmentation color (representing car ID) is not uniquely tied to a single bounding box.

A potential solution is to modify the evaluation logic: if two boxes share the same color, the system should count those points only for the box that contains the majority of that color’s points. In this way we can encounter this problem. see image (43) and (44)

4.2.2 Missing Ground Truth Bounding Boxes

In this scenes [Image ids:360(9), 1461(25), 1791 (29), 2712(38)], the segmented cars were visible and colored, but corresponding ground truth bounding boxes were absent. As a result, no precision or recall was calculated for those instances.

4.2.3 Too many GT boxes on one car

In the image Id 2449(37), we observed one problem, which was the presence of multiple ground truth (GT) bounding boxes overlapping a single physical car. As a result, the

model interpreted this as multiple distinct cars and generated several detections in the same area. However, these detections corresponded to only one actual vehicle. we have solve this problem with the help of above proposed solution (in section 4.2.1).




Figure 43: Too many BB on one car

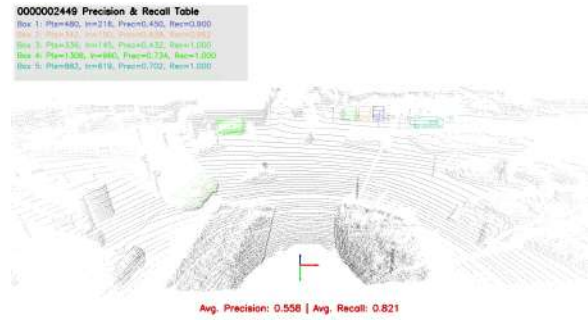


Figure 44: updated result

5 Conclusion

The proposed solution demonstrates promising results; however, several limitations were identified, as detailed in the image-specific error section. These errors impact the overall reliability of the solution. with further work and error correction, the method has potential for real-time applications. However, at its current stage, the system is not yet reliable enough for deployment in real-world automotive environments.

6 References

1. YOLOv8x-seg "Accessed June 13, 2025"
<https://docs.ultralytics.com/models/yolov8/overview>
2. KITTI 360 : <https://www.cvlibs.net/datasets/kitti-360/index.php>
3. KITTI 360 SUB Data set <https://elearning.rwu.de/mod/resource/view.php?id=213354>